

IE406

Machine Learning

Lab 1

Group 28

201801015 - Shantanu Tyagi

201801076 - Shivani Nandani

201801407 - Pratvi Shah

201801408 - Arkaprabha Banerjee

September 4, 2021

```
[143]: import numpy as np
import matplotlib.pyplot as plt
from mpl_toolkits import mplot3d
import pandas as pd
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import normalize
from yellowbrick.regressor import ResidualsPlot
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import MinMaxScaler
```

1 Question 1

```
[93]: def L(theta):
return theta*theta


theta = np.arange(-10,10,0.1)
L_theta = []
for t in theta:
L_theta.append(L(t))

min_L_theta = min(L_theta)
min_loc = np.argwhere(L_theta == min_L_theta)[0,0]

print('Minimum value of L() is', round(min_L_theta), ' at = ',
→round(min_L_theta))
plt.figure()
plt.plot(theta, L_theta, 'k-.', linewidth=3)
plt.plot(theta[min_loc], min_L_theta, 'ro', markersize=10)
```

```
plt.show()
```

Minimum value of $L()$ is 0.0 at $= 0.0$



2 Question 2

```
[108]: def L(theta1, theta2):
        return theta1*theta1 + theta2*theta2

        theta1 = np.arange(-10,10,0.1)
        theta2 = np.arange(-10,10,0.1)
        L_theta1_theta2 = np.zeros((len(theta1),len(theta2)))
        for i in range(len(theta1)):
            for j in range(len(theta2)):
                L_theta1_theta2[i][j] = L(theta1[i], theta2[j])

        min_L_theta1_theta2 = np.min(L_theta1_theta2)
        print('Minimum value of L(1, 2) is', round(min_L_theta1_theta2))
        min_loc = np.argwhere(L_theta1_theta2 == min_L_theta1_theta2)


        x, y = np.meshgrid(theta1, theta2)
        z = L(x, y)
```

```

fig = plt.figure(figsize=(10, 10))
ax = plt.axes(projection='3d')
ax.plot_wireframe(x, y, z, color='blue', linewidth=0.2)
ax.plot(theta1[min_loc[0,0]], theta2[min_loc[0,1]], z
→z[min_loc[0,0],min_loc[0,1]], 'ro', markersize=8)
ax.set_xlabel('1')
ax.set_ylabel('2')
ax.set_zlabel('L(1, 2)')
plt.show()

```

Minimum value of $L(1, 2)$ is 0.0



The minimum value of L is obtained at θ

2.1 Part A

```
[117]: def L(x, y, theta0, theta1):
        return np.sum(np.square(y - (theta0 + (theta1*x))))

        theta0 = np.arange(-50,50,0.05)
        theta1 = np.arange(-1,1,0.001)
        L_theta0_theta1 = np.zeros((len(theta1),len(theta0)))
        data = pd.read_excel('data.xlsx')
        data_x = np.reshape(np.array(data.x), (-1, 1))
        data_y = np.reshape(np.array(data.y), (-1, 1))

        for i in range(len(theta1)):
            for j in range(len(theta0)):
                L_theta0_theta1[i][j] = L(data_x, data_y, theta0[j], theta1[i])

        min_L_theta0_theta1 = np.min(L_theta0_theta1)
        print('Minimum value of L(0, 1) is', min_L_theta0_theta1)

        min_loc = np.argwhere(L_theta0_theta1 == min_L_theta0_theta1)
        print('0 and 1 values for minimum value of L(0, 1) are ',
        →theta0[min_loc[0,1]], theta1[min_loc[0,0]], ' respectively.')


        x, y = np.meshgrid(theta0, theta1)
        z = L_theta0_theta1

        fig = plt.figure(figsize=(10, 10))
        ax = plt.axes(projection='3d')
        ax.plot_wireframe(x, y, z, color='blue', linewidth=0.2)
        ax.plot(theta0[min_loc[0,1]], theta1[min_loc[0,0]],
        →z[min_loc[0,0],min_loc[0,1]], 'ro', markersize = 8)
        ax.set_xlabel('0')
        ax.set_ylabel('1')
        ax.set_zlabel('L(0, 1)')
        plt.show()
```

Minimum value of L(0, 1) is 1595.086384000078

0 and 1 values for minimum value of L(0, 1) are 47.399999999999446

-0.007999999999999119 respectively.



2.2 Part B

```
[118]: #Question 3b
temp = np.copy(data_x)
x = np.ones((np.size(temp,0),2))
x[:,1] = temp[:,0]

print('determinant of x.T*x = ',np.linalg.det(np.matmul(x.T,x)))

theta = np.matmul(np.linalg.inv(np.matmul(x.T,x)),np.matmul(x.T,data_y))
print('For minima, theta0 = ',theta[0,0],' and theta1 = ',theta[1,0])
```

determinant of x.T*x = 5618077959.999997

For minima, theta0 = 49.23762989433493 and theta1 = -0.008611934783475328

Using θ

```
[123]: #Question 4
        LS = np.sum(np.square(data_y - (theta[0,0] + (theta[1,0]*temp))))
        print('Value of L using theta values from LS method(Pseudo Inverse) =')
        print(LS)
        print()
        for i in range(10):
            L_other_value = np.sum(np.square(data_y - (np.random.randint(theta0.
            shape[0]) - (np.random.randint(theta1.shape[0])*temp))))
            print('Value of L for some other theta = ',L_other_value)
```

```
Value of L using theta values from LS method(Pseudo Inverse) =
1572.6503668922921
```

```
Value of L for some other theta = 2008722638480283.8
Value of L for some other theta = 1119326591644803.8
Value of L for some other theta = 256774493245402.75
Value of L for some other theta = 227493761746521.75
Value of L for some other theta = 2458405151084183.0
Value of L for some other theta = 201287534533077.75
Value of L for some other theta = 148445180762662.75
Value of L for some other theta = 2027591949333187.8
Value of L for some other theta = 795214078513074.8
Value of L for some other theta = 21985849704444.75
```

We see that the value L for the θ

2.3 Part A

```
[127]: x = np.array([
        [1, 2],
        [2, 4],
        [3, 6],
        [4, 8]
    ])
        y = np.array([
        [2],
        [3],
        [4],
        [5]
    ])

        print(' X Shape ->', x.shape, '\n', 'Y Shape ->', y.shape)
```

```
X Shape -> (4, 2)
Y Shape -> (4, 1)
```


```
[128]: reg = LinearRegression().fit(x, y)
print('Score: ', reg.score(x, y))
print('Coefficient: ', reg.coef_)
print('Intercept: ', reg.intercept_)
```


```
Score: 1.0
Coefficient: [[0.2 0.4]]
Intercept: [1.]
```

```
[130]: fig = plt.figure()
ax = plt.axes(projection='3d')
ax.plot_wireframe(x[:,0], x[:,1], y, color='k', linewidth=0.7,
→label='actual')
ax.plot_wireframe(x[:,0], x[:,1], y_pred, color='r', linewidth=0.7,
→label='predicted')
ax.set_xlabel('x[:,0]')
ax.set_ylabel('x[:,1]')
ax.set_zlabel('y')
ax.view_init(90, 0)
plt.legend()
plt.title('Top View')
plt.show()

fig = plt.figure()
ax = plt.axes(projection='3d')
ax.plot_wireframe(x[:,0], x[:,1], y, color='k', linewidth=0.7,
→label='actual')
ax.plot_wireframe(x[:,0], x[:,1], y_pred, color='r', linewidth=0.7,
→label='predicted')
ax.set_xlabel('x[:,0]')
ax.set_ylabel('x[:,1]')
ax.set_zlabel('y')
plt.legend()
plt.title('Side View')
plt.show()
```

Top View





2.4 Part B

Since $(X^T X)$ is non-invertible, therefore we obtain an error when solving via normal equations method

```
[136]: def normal(X,y):
        return np.dot(np.dot(np.linalg.inv(np.dot(X.T,X)),X.T),y)

        normal(x,y)
```

```
-----
LinAlgError                                Traceback (most recent call last)
<ipython-input-136-ccba76942472> in <module>
      2     return np.dot(np.dot(np.linalg.inv(np.dot(X.T,X)),X.T),y)
      3
----> 4 normal(x,y)

<ipython-input-136-ccba76942472> in normal(X, y)
      1 def normal(X,y):
```

```

----> 2     return np.dot(np.dot(np.linalg.inv(np.dot(X.T,X)),X.T),y)
      3
      4 normal(x,y)

<__array_function__ internals> in inv(*args, **kwargs)

c:
→\users\geeksa67\appdata\local\programs\python\python38\lib\site-packages\numpy linalg
    \linalg.py in inv(a)
      545     signature = 'D->D' if isComplexType(t)
    else 'd->d'
      546     extobj = get_linalg_error_extobj(_raise_linalgerror_singular)
--> 547     ainv = _umath_linalg.inv(a,
signature=signature, extobj=extobj)
      548     return wrap(ainv.astype(result_t, copy=
False))
      549

c:
→\users\geeksa67\appdata\local\programs\python\python38\lib\site-packages\numpy linalg\linalg
→py in
    _raise_linalgerror_singular(err, flag)
      95
      96 def _raise_linalgerror_singular(err, flag):
--> 97     raise LinAlgError("Singular matrix")
      98
      99 def _raise_linalgerror_nonposdef(err, flag):

LinAlgError: Singular matrix

```

```

[142]: x_trans = np.transpose(x)
      XX = np.dot(x_trans, x)
      XY = np.dot(x_trans, y)

      print('Checking invertability')
      print('-----')

      print('X\'X\n', XX)
      print('Rank of X\'X -> ', np.linalg.matrix_rank(XX))

      print('-----')
      print('Since the matrix is rank deficient, thus matrix X\'X is not_
→invertible.')

```

```

Checking invertability
-----
X'X

```

```
[[ 30  60]
 [ 60 120]]
Rank of X'X -> 1
```

Since the matrix is rank deficient, thus matrix $X'X$ is not invertible.

The Inverse of ATA doesn't exist in the above questions hence the approach of normal equations wont work here. This is primarily because ATA is a rank-deficient matrix hence the inverse doesn't exist.

The LinearRegression function is a wrapper function for `scipy.linalg.lstsq` and other associated methods. For rank deficient matrices where the inverse of a matrix does not exist by traditional methods (normal equations), these libraries compute a least-squares solution to equation $Ax = b$ by Computing a vector x such that the 2-norm $|b - Ax|$ is minimized. This is achieved by using LAPACK (Linear Algebra Package by NETLIB) which in turn uses SVD or QR (Orthogonal Factorization) based techniques to calculate the inverse.

3 Question 6

3.1 Part A

```
[386]: real_estate_data = pd.read_excel('real_estate_valuation_dataset.xlsx')
       real_estate_data
```

```
[386]:
```

No	X1	transaction date	X2 house age	\
0	1	2012.916667		32.0
1	2	2012.916667		19.5
2	3	2013.583333		13.3
3	4	2013.500000		13.3
4	5	2012.833333		5.0
..
409	410	2013.000000		13.7
410	411	2012.666667		5.6
411	412	2013.250000		18.8
412	413	2013.000000		8.1
413	414	2013.500000		6.5

	X3 distance to the nearest MRT station	X4 number of convenience stores
→ \		
0		84.87882
→ 10		
1		306.59470
→ 9		
2		561.98450
→ 5		
3		561.98450
→ 5		

```

4          390.56840
→ 5
..          ...
→...
409          4082.01500
→ 0
410          90.45606
→ 9
411          390.96960
→ 7
412          104.81010
→ 5
413          90.45606
→ 9

```

	X5 latitude	X6 longitude	Y house price of unit area
0	24.98298	121.54024	37.9
1	24.98034	121.53951	42.2
2	24.98746	121.54391	47.3
3	24.98746	121.54391	54.8
4	24.97937	121.54245	43.1
..
409	24.94155	121.50381	15.4
410	24.97433	121.54310	50.0
411	24.97923	121.53986	40.6
412	24.96674	121.54067	52.5
413	24.97433	121.54310	63.9

[414 rows x 8 columns]

```

[387]: x1 = np.reshape(np.array(real_estate_data['X1 transaction date']), (-1,
→1))
x2 = np.reshape(np.array(real_estate_data['X2 house age']), (-1, 1))
x3 = np.reshape(np.array(real_estate_data['X3 distance to the nearest
→MRT station']), (-1, 1))
x4 = np.reshape(np.array(real_estate_data['X4 number of convenience
→stores']), (-1, 1))
x5 = np.reshape(np.array(real_estate_data['X5 latitude']), (-1, 1))
x6 = np.reshape(np.array(real_estate_data['X6 longitude']), (-1, 1))
X = np.hstack((x1, x2, x3, x4, x5, x6))
y = np.reshape(np.array(real_estate_data['Y house price of unit area']),
→(-1, 1))

```

```

[388]: reg = LinearRegression().fit(X, y)
print('Score: ', reg.score(X, y))
print('Coefficient: ', reg.coef_)

```

```
print('Intercept: ', reg.intercept_)

y_pred1 = reg.predict(X)

print('RMSE ->', np.sqrt(mean_squared_error(y, y_pred1)))
```

```
Score: 0.5823850447850787
Coefficient: [[ 5.14901721e+00 -2.69696735e-01 -4.48750825e-03  1.13332498e+00
                2.25470143e+02 -1.24290612e+01]]
Intercept: [-14441.98271918]
RMSE -> 8.782312975361124
```

3.2 Part B

The sign of a regression coefficient tells you whether there is a positive or negative correlation between each independent variable and the dependent variable. A positive coefficient indicates that as the value of the independent variable increases, the mean of the dependent variable also tends to increase. A negative coefficient suggests that as the independent variable increases, the dependent variable tends to decrease.

The coefficient value signifies how much the mean of the dependent variable changes given a one-unit shift in the independent variable while holding other variables in the model constant. This property of holding the other variables constant is crucial because it allows you to assess the effect of each variable in isolation from the others.

We can fit a LinearRegression model on the regression dataset and retrieve the `coeff_` property that contains the coefficients found for each input variable. Since the X values are not standardized, we cannot comment on the importance of the feature score.

3.3 Part C

```
[389]: X_norm = (X - X.min())/(X.max()-X.min())
reg = LinearRegression().fit(X_norm, y)
print('Score: ', reg.score(X_norm, y))
print('Coefficient: ', reg.coef_)
print('Intercept: ', reg.intercept_)

y_pred2 = reg.predict(X_norm)


print('RMSE ->', np.sqrt(mean_squared_error(y, y_pred2)))
```

```
Score: 0.5823850447850758
Coefficient: [[ 3.34069318e+04 -1.74979808e+03 -2.91150478e+01  7.35303628e+03
                1.46285502e+06 -8.06400099e+04]]
Intercept: [-14441.98271918]
RMSE -> 8.782312975361156
```

```
[390]: fig = plt.figure(figsize = (10, 5))
```

```
# creating the bar plot
plt.bar([1,2,3,4,5,6], np.log10(abs(reg.coef_[0])), color = 'maroon',
width = 0.4)

plt.xlabel("Feature")
plt.ylabel("Score(log10)")
plt.title("Score for each feature")
#plt.yscale('log')
plt.show()
```

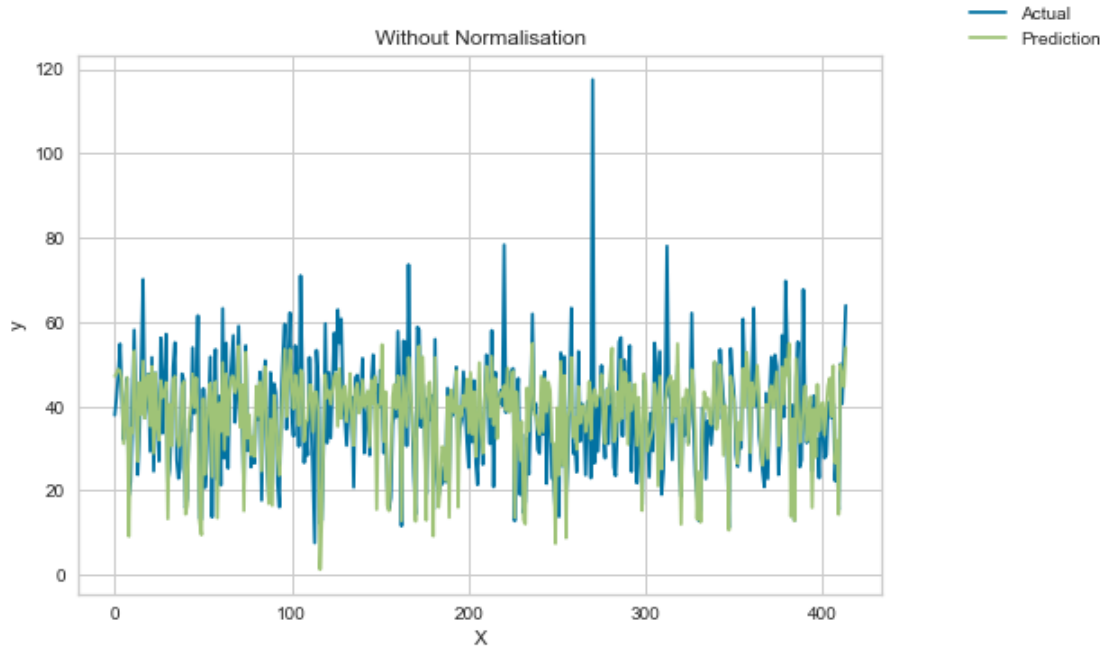


The sign of a regression coefficient tells if there is a positive or negative correlation between each independent variable and dependent variable, i.e., if a feature is positively correlated, increase in its value increases the mean of the dependent and vice-versa. The value of the correlation coefficient gives an idea of how much a unit shift in the value of an independent variable affects the dependent variable (assuming all other variables to be constant). This gives an idea of the effect of each variable independent of the others.

A true sense of this can be achieved only after the normalization has been performed. Thus, once normalized, the coefficients can be used as a measure of the importance of each feature w.r.t. the dependent variable. Else, as mentioned above, the coefficients are not a good measure of the importance of the features. As seen in the graph, the 5th and 6th features are the most important ones, i.e., the latitude and longitude of the house matters.

```
[391]: #plt.plot(y)
plt.plot(y)
plt.plot(y_pred1)
plt.legend(["Actual", "Prediction"], loc=[1.1, 1])
```

```
plt.xlabel("X")
plt.ylabel("y")
plt.title("Without Normalisation")
plt.show()
```




3.4 Part D

```
[392]: # Create the train and test data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
→random_state=42)

# Instantiate the linear model and visualizer
model = LinearRegression()
visualizer = ResidualsPlot(model)


visualizer.fit(X_train, y_train) # Fit the training data to the
→visualizer

visualizer.score(X_test, y_test) # Evaluate the model on the test data
visualizer.show() # Finalize and render the figure
plt.show()
```



3.5 Part E

```
[400]: import seaborn as sns
plt.figure(figsize=(12,10))
X_norm = pd.DataFrame(X_norm)
X_norm.columns = [1,2,3,4,5,6]
cor = X_norm.corr()
sns.heatmap(cor, annot=True, cmap=plt.cm.CMRmap_r)
plt.show()
```

```
[401]: # with the following function we can select highly correlated features
# it will remove the first feature that is correlated with anything
→other feature
def correlation(dataset, threshold):
    col_corr = set() # Set of all the names of correlated columns
    corr_matrix = dataset.corr()
    for i in range(len(corr_matrix.columns)):
        for j in range(i):
            if abs(corr_matrix.iloc[i, j]) > threshold: # we are interested in
→absolute coeff value
                colname = corr_matrix.columns[i] # getting the name of column
                col_corr.add(colname)
    return col_corr
```

```
[409]: rmse = [0 for i in range(10)]
for i in range(1,11):
```

```


corr_features = correlation(X_norm,i/10)
if len(set(corr_features)) != 6:
    temp = X_norm.drop(corr_features,axis=1)
    reg = LinearRegression().fit(np.array(temp), y)
    y_pred2 = reg.predict(np.array(temp))
    rmse[i-1] = np.sqrt(mean_squared_error(y, y_pred2))

```

```

[424]: plt.plot([i/10 for i in range(1,11)],rmse)
plt.xlabel("Threshhold")
plt.ylabel("RMSE")
plt.show()

```



```

[416]: corr_features = correlation(X_norm,0.7)
corr_features

```

```

[416]: {6}

```

One feature is removed giving the no of features retained = 5 out of 6 features retained are columns = {1, 2, 3, 4, 5} with 6th column getting removed